

```
In [1]: import keras
        from keras.datasets import fashion_mnist
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
```

```
In [2]: # Load Data
        (X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()

        Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
        29515/29515 [=====] - 0s 4us/step
        Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
        26421880/26421880 [=====] - 38s 1us/step
        Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
        5148/5148 [=====] - 0s 0s/step
        Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
        4422102/4422102 [=====] - 7s 1us/step
```

```
In [3]: # Preprocessing
        X_train = X_train.reshape(60000, 28, 28, 1).astype('float32')
        X_test = X_test.reshape(10000, 28, 28, 1).astype('float32')

        X_train /= 255
        X_test /= 255

        n_classes = 10
        y_train = keras.utils.to_categorical(y_train, n_classes)
        y_test = keras.utils.to_categorical(y_test, n_classes)
```

```
In [4]: # Design neural network architecture
model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(n_classes, activation='softmax'))
```

```
In [5]: # Model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
conv2d_1 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 12, 12, 64)	0
dropout (Dropout)	(None, 12, 12, 64)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 128)	1179776
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290

=====
Total params: 1,199,882
Trainable params: 1,199,882
Non-trainable params: 0

```
In [6]: # Configure the model
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [7]: # Train the model
model.fit(X_train, y_train, batch_size=128, epochs=10, verbose=1, validation_data=(X_test, y_test))
```

```
Epoch 1/10
469/469 [=====] - 96s 200ms/step - loss: 0.5164 - accuracy: 0.8167 - val_loss: 0.3439 - val_accuracy: 0.8759
Epoch 2/10
469/469 [=====] - 94s 201ms/step - loss: 0.3393 - accuracy: 0.8787 - val_loss: 0.2844 - val_accuracy: 0.8964
Epoch 3/10
469/469 [=====] - 86s 183ms/step - loss: 0.2920 - accuracy: 0.8943 - val_loss: 0.2706 - val_accuracy: 0.8999
Epoch 4/10
469/469 [=====] - 81s 173ms/step - loss: 0.2604 - accuracy: 0.9054 - val_loss: 0.2453 - val_accuracy: 0.9094
Epoch 5/10
469/469 [=====] - 99s 211ms/step - loss: 0.2350 - accuracy: 0.9145 - val_loss: 0.2339 - val_accuracy: 0.9124
Epoch 6/10
469/469 [=====] - 94s 201ms/step - loss: 0.2171 - accuracy: 0.9195 - val_loss: 0.2312 - val_accuracy: 0.9145
Epoch 7/10
469/469 [=====] - 92s 196ms/step - loss: 0.2006 - accuracy: 0.9251 - val_loss: 0.2168 - val_accuracy: 0.9228
Epoch 8/10
469/469 [=====] - 92s 196ms/step - loss: 0.1836 - accuracy: 0.9314 - val_loss: 0.2206 - val_accuracy: 0.9209
```

```
In [8]: # Evaluate the model
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Test loss: 0.20950645208358765
Test accuracy: 0.9273999929428101
```